

# Lab: Using Counters and a Map-Only Job

In this lab you will create a Map-only MapReduce job.

Your application will process a web server's access log to count the number of times gifs, jpegs, and other resources have been retrieved. Your job will report three figures: number of gif requests, number of jpeg requests, and number of other requests.

## Instructions

### 1. Input Data

Unpack the `access_log.gz` file and upload it to HDFS. You can unpack the file with the following command:

```
gunzip access_log.gz
```

### 2. Source code

You will find the source files in the `src/stubs` directory. You will also find a `src/hints` directory that contains a copy of the source files modified to make the assignment easier to complete, and a `src/solution` directory that contains a copy of the source files modified to complete the assignment.

### 3. Complete the `ImageCounter` and `ImageCounterMapper` classes in the `src/stubs` directory to count the number of occurrences of files of type GIF, JPEG, and "other". Use counters to return the counts from the map task.

### 4. Assemble and test your solution.

## Hints

1. You should use a Map-only MapReduce job, by setting the number of Reducers to 0 in the driver code.
2. You should test your code against the smaller version of the access log before testing against the full log file.
3. Use a counter group such as `ImageCounter`, with names `gif`, `jpeg`, and `other`.
4. In your driver code, retrieve the values of the counters after the job has completed and report them using `System.out.println`.
5. The output folder on HDFS will contain Mapper output files which are empty, because the Mappers did not write any data.